

```

import java.io.IOException;

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Put;
import org.apache.hadoop.hbase.util.Bytes;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.mapred.lib.NullOutputFormat;
import org.apache.hadoop.util.*;

/*
 * This importer is used to read the raw input data on HDFS and write it to HBase
 * prior to running this, create the appropriate tables in hbase (via hbase shell)
 * create 'stations',{NAME => 'info', VERSIONS => 1}
 * create 'observations, {NAME => 'data', VERSIONS => 1}
 * for this case, we are only interested in the latest version, hence VERSIONS is set to 1.
 */
public class HBaseTemperatureImporter extends Configured implements Tool {

    // Inner-class for map
    static class HBaseTemperatureMapper<K, V> extends MapReduceBase implements
        Mapper<LongWritable, Text, K, V> {
        private NcdcRecordParser parser = new NcdcRecordParser();
        private HTable table;

        public void map(LongWritable key, Text value,
            OutputCollector<K, V> output, Reporter reporter)
            throws IOException {
            // input key-value as TextInputFormat (default), with the byte offset and a long string of
            integers containing data about
            // weather station, air temperature etc.
            // using the parser to parse the string and determine if temperature is value
            parser.parse(value.toString());
            if (parser.isValidTemperature()) {
                // adds valid temperatures to the observations HBase table into the data:air-temp
                column
                // row key generated via custom function as stationID+dateInReverseEpoch
                byte[] rowKey = RowKeyConverter.makeObservationRowKey(parser.getStationId(),
                    parser.getObservationDate().getTime());
                Put p = new Put(rowKey);
                p.add(HBaseTemperatureCli.DATA_COLUMNFAMILY,
                    HBaseTemperatureCli.AIRTEMP_QUALIFIER,
                    Bytes.toBytes(parser.getAirTemperature()));
                table.put(p);
            }
        }
    }

    /*

```

* Creates a HTable instance once against the observations table and use it afterward in
map invocations talking to HBase

```
        */  
    public void configure(JobConf jc) {  
        super.configure(jc);  
        // Create the HBase table client once up-front and keep it around  
        // rather than create on each map invocation.  
        try {  
            this.table = new HTable(new HBaseConfiguration(jc), "observations");  
        } catch (IOException e) {  
            throw new RuntimeException("Failed HTable construction", e);  
        }  
    }  
  
    @Override  
    public void close() throws IOException {  
        super.close();  
        table.close();  
    }  
}  
  
public int run(String[] args) throws IOException {  
    if (args.length != 1) {  
        System.err.println("Usage: HBaseTemperatureImporter <input>");  
        return -1;  
    }  
    JobConf jc = new JobConf(getConf(), getClass());  
    FileInputFormat.addInputPath(jc, new Path(args[0]));  
    jc.setMapperClass(HBaseTemperatureMapper.class);  
    jc.setNumReduceTasks(0);  
    jc.setOutputFormat(NullOutputFormat.class);  
    JobClient.runJob(jc);  
    return 0;  
}  
  
public static void main(String[] args) throws Exception {  
    int exitCode = ToolRunner.run(new HBaseConfiguration(),  
        new HBaseTemperatureImporter(), args);  
    System.exit(exitCode);  
}  
}
```