

```

import java.io.IOException;

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Put;
import org.apache.hadoop.hbase.util.Bytes;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.mapred.lib.NullOutputFormat;
import org.apache.hadoop.util.*;

/**
 * A MapReduce application to count the number of rows in an HBase table
 */
public class RowCounter extends Configured implements Tool {
    // Name of this program
    static final String NAME = "rowcounter";

    /* HBase TableMap interface is a specialization of org.apache.hadoop.mapred.Mapper
     * that sets the map input types passed by TableInputFormat (set by JobConf below via
     * TableMapUtil.initTableMapJob() utility method
     * this mapper checks all the columns of a row, if all are empty, it doesn't count the row,
     * otherwise it increments Counters.ROWS by one
     */
    static class RowCounterMapper implements TableMap<ImmutableBytesWritable,
    RowResult> {
        private static enum Counters {ROWS}

        public void map(ImmutableBytesWritable row, RowResult value,
        OutputCollector<ImmutableBytesWritable, RowResult> output,
        Reporter reporter) throws IOException {
            boolean content = false;
            for (Map.Entry<byte [], Cell> e: value.entrySet()){
                Cell cell = e.getValue();
                if (cell != null && cell.getValue().length > 0) {
                    content = true;
                    break;
                }
            }
            if (!content) {
                // Don't count rows that are all empty values.
                return;
            }
            // Give out some value every time. We are only interested in the row/key
            reporter.incrCounter(Counters.ROWS, 1);
        }

        public void configure (JobConf jc) {
            // Nothing to do.

```

```

    }

    public void close() throws IOException {
        // Nothing to do.
    }
}

/*
 * This method parses arguments added to the configuration that were passed on the
command line figuring the table and columns
 * to run RowCounter against. It also invokes TableMapUtil.initTableMapJob() utility
method,
 * which among other things such as setting the map class to use, sets the input format
to TableInputFormat.
 */
public JobConf createSubmittableJob(String[] args) throws IOException {
    JobConf c = new JobConf(getConf(), getClass());
    c.setJobName(NAME);
    // Columns are space delimited
    StringBuilder sb = new StringBuilder();
    final int column offset = 2;
    for (int i = columnoffset; i < args.length, i++) {
        if (i > columnoffset) {
            sb.append(" ");
        }
        sb.append(args[i]);
    }
    // Second argument is the table name.
    TableMapReduceUtil.initTableMapJob(args[1], sb.toString(),
RowCounterMapper.class, ImmutableBytesWritable.class,
    RowResult.class, c);
    c.setNumReduceTasks(0);
    // First arg is the output directory.
    FileOutputFormat.setOutputPath(c, new Path(args[0]));
    return c;
}

static int printUsage() {
    System.out.println(NAME + " <outputdir> <tablename> <column1>
[<column2>...]");
    return -1;
}

/**
 * run() method belongs to Tool implementation
 */
public int run(final String[] args) throws Exception {
    // Make sure there are at least 3 parameters
    if (args.length < 3) {
        System.err.println("ERROR: Wrong number of parameters: " +
args.length);
    }
}

```

```

        return printUsage();
    }
    JobClient.runJob(createSubmittableJob(args));
    return 0;
}

/**
 * RowCounter's main() method does not invoke its own run() method directly. Instead it
calls ToolRunner's static run() method,
 * which takes care of creating a Configuration object for the Tool, before calling its run()
method. ToolRunner also uses a
 * GenericOptionsParser to pick up any standard options specified on the command line,
and set them on the Configuration instance.
 */
public static void main(String[] args) throws Exception {
    HBaseConfiguration c = new HBaseConfiguration();
    int errCode = ToolRunner.run(c, new RowCounter(), args);
    System.exit(errCode);
}
}

```